

Harrison Li

Berkeley, CA | liharrison@berkeley.edu | 425 698 7779 | LinkedIn | GitHub

Education

University of California, Berkeley, BA in Data Science Aug 2025 – Present

- **Coursework:** Program & Data Structures, Data Techniques, Linear Algebra, Discrete Math

Prestige Stem Academy, Issaquah High School

Sep 2021 – Jun 2025

- **Coursework:** AP CSA, AP Stats, Advanced CS Topics & Proj.

Skills

- **Languages:** Python, JavaScript, TypeScript, Java, HTML/CSS, SQL
- **ML / AI:** PyTorch, TensorFlow, LangChain, LangGraph, MCP, LSTMs, Hugging Face, OpenCV, RAG
- **Full-Stack:** React, Next.js, Tailwind CSS, Framer Motion, Supabase, Git

Experience

Intern, Quo AI & Starflux AI Dec 2025 – Present

- **Quo AI:** Developed and deployed a full-stack web platform; building agentic automation systems for home buying and selling workflows.
- **Starflux AI:** Researching reliable agent behavior; synthesis, simulation, and evaluation of agents for multi-agent and reinforcement learning environments.

Undergraduate Researcher, Berkeley Artificial Intelligence Research (BAIR) Oct 2025 – Present

- Collaborating with UCSF Neuroscience and licensed Speech-Language Pathologists to develop automated clinical decision-support systems for speech and language disorder diagnosis.
- Designing simulation environments modeling speech-language pathology workflows as POMDPs to study information acquisition, risk-sensitive decision-making, and diagnostic safety.
- Built state-of-the-art Primary Progressive Aphasia dysfluency modeling system with clinical grounding.

ML Researcher, Algoverse May 2025 – Oct 2025

- Designed and implemented evaluation framework for LLM agentic reasoning on Tower of Hanoi.
- Developed novel experimental methodology to measure policy divergence in multi-step agentic search environments.
- Co-authored paper accepted to FoRLM @ NeurIPS 2025 on limitations of emergent reasoning in LLM-based agentic frameworks.

Projects

Santé (TreeHacks 2026) Dec 2025 – Present

- Built a multi-agent voice intelligence platform that extracts acoustic biomarkers and phoneme-level disfluency signals from real-time patient calls to generate clinician-facing diagnostic reports.
- Architected a full-stack pipeline with Next.js, FastAPI, OpenAI Realtime, and serverless GPU inference, parallelizing live transcription, openSMILE acoustic analysis, and DysfluentWFST phoneme extraction.

BeWear (CalHacks 2025) Oct 2025

- Built a full-stack system for visual clothing brand identification and ethical scoring across sustainability, labor, and transparency metrics.
- Implemented an agentic pipeline for real-time greenwashing detection, integrating web search, Elasticsearch, and state-managed frontend updates.

DoodleAI Feb 2025 – May 2025

- Built a deep learning system to classify and generate sketches from the QuickDraw dataset, including a CNN for 50-category classification and a custom LSTM-based sketch generator.
- Developed a Tkinter drawing interface with real-time CNN predictions.